

UNIVERSITÀ
DEGLI STUDI
DI PADOVA



DIPARTIMENTO
DI INGEGNERIA
DELL'INFORMAZIONE



IAS-Lab

Intelligent Autonomous
Systems Laboratory

3D Data Processing Ceres-Solver Tutorial

Alberto Pretto

Goal

Solve robustified **non-linear least squares** problems of the form

The diagram shows the optimization problem
$$\min_{\mathbf{x}} \quad \frac{1}{2} \sum_i \rho_i \left(\|f_i(x_{i_1}, \dots, x_{i_k})\|^2 \right)$$
 with several annotations: a red oval labeled "Residual Block" encircles the entire sum; a blue oval labeled "Parameter Block/s" encircles the input parameters x_{i_1}, \dots, x_{i_k} ; a blue line points from the label "Loss Function (Huber, Cauchy ...)" to the ρ_i term; and a blue line points from the label "Cost Function" to the $\|f_i\|^2$ term.

$$\min_{\mathbf{x}} \quad \frac{1}{2} \sum_i \rho_i \left(\|f_i(x_{i_1}, \dots, x_{i_k})\|^2 \right)$$

Residual Block

Loss Function (Huber, Cauchy ...)

Cost Function

Parameter Block/s

Example

Find the minimum of the function

$$\frac{1}{2}(10 - x)^2$$

Solve it with Ceres:

- 1) Write a functor that will evaluate the residual
- 2) Build the non-linear least squares problem
- 3) Setup and run the solver

```
struct CostFunction {  
    template <typename T>  
    bool operator()(const T* const x, T* residual) const {  
        residual[0] = 10.0 - x[0];  
        return true;  
    }  
};
```

```
int main(int argc, char** argv) {  
    google::InitGoogleLogging(argv[0]);  
  
    // The variable to solve for with its initial value.  
    double initial_x = 5.0;  
    double x = initial_x;  
  
    // Build the problem.  
    Problem problem;  
  
    // Set up the only cost function (also known as residual). This uses  
    // auto-differentiation to obtain the derivative (jacobian).  
    CostFunction* cost_function =  
        new AutoDiffCostFunction<CostFunction, 1, 1>(new CostFunction);  
    problem.AddResidualBlock(cost_function, nullptr, &x);  
  
    // Run the solver!  
    Solver::Options options;  
    options.linear_solver_type = ceres::DENSE_QR;  
    options.minimizer_progress_to_stdout = true;  
    Solver::Summary summary;  
    Solve(options, &problem, &summary);  
  
    std::cout << summary.BriefReport() << "\n";  
    std::cout << "x : " << initial_x  
        << " -> " << x << "\n";  
  
    return 0;  
}
```

Example

Find the minimum of the function

$$\frac{1}{2}(10 - x)^2$$

OUTPUT

iter	cost	cost_change	gradient	step	tr_ratio	tr_radius	ls_iter	iter_time	total_time
0	4.512500e+01	0.00e+00	9.50e+00	0.00e+00	0.00e+00	1.00e+04	0	5.33e-04	3.46e-03
1	4.511598e-07	4.51e+01	9.50e-04	9.50e+00	1.00e+00	3.00e+04	1	5.00e-04	4.05e-03
2	5.012552e-16	4.51e-07	3.17e-08	9.50e-04	1.00e+00	9.00e+04	1	1.60e-05	4.09e-03

Ceres Solver Report: Iterations: 2, Initial cost: 4.512500e+01, Final cost: 5.012552e-16, Termination: CONVERGENCE
x : 5.0 -> 10

Automatic differentiation

- Ceres can compute automatically the derivatives wrt the parameters vector while computing residuals
- The parameters can be divided into "blocks", as for example done in bundle adjustment ("camera" blocks and "point" blocks), for simplify managing the sparsity

Adding residuals

For each residual, we need to add a corresponding "residual block" to the optimization problem:

```
ceres::Problem problem;
for( /* iterate for each data point */ )
{
    ceres::CostFunction* cost_function = ...;

    problem.AddResidualBlock( cost_function,
                             param_block1, param_block2, ...);
}
```

Adding residuals

We need to define a functor (just a class or struct which defines the operator()) that computes the residual:

```
struct Functor
{
    template <typename T> bool operator() (const T* const param_block1,
                                           const T* const param_block2,
                                           ...,
                                           T* residuals) const
    {
        // Compute the residuals given the input parameters blocks
        return true; // Success
    }
}
```

Adding residuals

Then we use this functor to construct the cost function:

```
Functor funct = new Functor(...);  
ceres::CostFunction* cost_function = new  
ceres::AutoDiffCostFunction<Functor, Nr, Nb1, Nb2, ... >(funct);
```

N_r: dimension of a single residual

N_{b1}: dimension of parameters block 1

N_{b2}: dimension of parameters block 2


....

Curve Fitting

Given a set of observed data points, find the best fitting exponential curve

$$y = e^{mx+c}$$

```
struct ExponentialResidual {  
    ExponentialResidual(double x, double y)  
        : x_(x), y_(y) {}  
  
    template <typename T>  
    bool operator()(const T* const m, const T* const c, T* residual) const {  
        residual[0] = y_ - exp(m[0] * x_ + c[0]);  
        return true;  
    }  
  
private:  
    // Observations for a sample.  
    const double x_;  
    const double y_;  
};
```




```
double m = 0.0;  
double c = 0.0;  
  
Problem problem;  
for (int i = 0; i < kNumObservations; ++i) {  
    CostFunction* cost_function =  
        new AutoDiffCostFunction<ExponentialResidual, 1, 1, 1>(  
            new ExponentialResidual(data[2 * i], data[2 * i + 1]));  
    problem.AddResidualBlock(cost_function, nullptr, &m, &c);  
}
```

Curve Fitting

Given a set of observed data points, find the best fitting exponential curve

$$y = e^{mx+c}$$



iter	cost	cost_change	gradient	step	tr_ratio	tr_radius	ls_iter	iter_time	total_time
0	1.211734e+02	0.00e+00	3.61e+02	0.00e+00	0.00e+00	1.00e+04	0	5.34e-04	2.56e-03
1	1.211734e+02	-2.21e+03	0.00e+00	7.52e-01	-1.87e+01	5.00e+03	1	4.29e-05	3.25e-03
2	1.211734e+02	-2.21e+03	0.00e+00	7.51e-01	-1.86e+01	1.25e+03	1	1.10e-05	3.28e-03
3	1.211734e+02	-2.19e+03	0.00e+00	7.48e-01	-1.85e+01	1.56e+02	1	1.41e-05	3.31e-03
4	1.211734e+02	-2.02e+03	0.00e+00	7.22e-01	-1.70e+01	9.77e+00	1	1.00e-05	3.34e-03
5	1.211734e+02	-7.34e+02	0.00e+00	5.78e-01	-6.32e+00	3.05e-01	1	1.00e-05	3.36e-03
6	3.306595e+01	8.81e+01	4.10e+02	3.18e-01	1.37e+00	9.16e-01	1	2.79e-05	3.41e-03
7	6.426770e+00	2.66e+01	1.81e+02	1.29e-01	1.10e+00	2.75e+00	1	2.10e-05	3.45e-03
8	3.344546e+00	3.08e+00	5.51e+01	3.05e-02	1.03e+00	8.24e+00	1	2.10e-05	3.48e-03
9	1.987485e+00	1.36e+00	2.33e+01	8.87e-02	9.94e-01	2.47e+01	1	2.10e-05	3.52e-03
10	1.211585e+00	7.76e-01	8.22e+00	1.05e-01	9.89e-01	7.42e+01	1	2.10e-05	3.56e-03
11	1.063265e+00	1.48e-01	1.44e+00	6.06e-02	9.97e-01	2.22e+02	1	2.60e-05	3.61e-03
12	1.056795e+00	6.47e-03	1.18e-01	1.47e-02	1.00e+00	6.67e+02	1	2.10e-05	3.64e-03
13	1.056751e+00	4.39e-05	3.79e-03	1.28e-03	1.00e+00	2.00e+03	1	2.10e-05	3.68e-03

Ceres Solver Report: Iterations: 13, Initial cost: 1.211734e+02, Final cost: 1.056751e+00, Termination: CONVERGENCE


Initial m: 0 c: 0

Final m: 0.291861 c: 0.131439


Robust Curve Fitting

$$y = e^{mx+c}$$

Without Loss Function



With Loss Function



Exploit loss functions for reducing the influence of outliers

```
problem.AddResidualBlock(cost_function, new CauchyLoss(0.5) , &m, &c);
```

References

- http://ceres-solver.org/nns_tutorial.html
- http://ceres-solver.org/nns_tutorial.html#bundle-adjustment